

# LLM-Guided PCA-Latent Corrosion Mask Forecasting

## Technical Report

(author name)

February 23, 2026

### Abstract

This report documents the mathematical foundations and software architecture of a corrosion forecasting pipeline designed to predict the temporal evolution of binary corrosion masks for Mg–4Ag biodegradable alloy wires. The masks are derived from in-situ synchrotron nano-CT imaging and captured at ten discrete degradation time-steps. The pipeline combines *Principal Component Analysis* (PCA) via randomised Singular Value Decomposition (SVD), a *Signed Distance Field* (SDF) representation,  $k$ -Nearest Neighbour ( $k$ NN) delta prediction in latent space, and a *Large Language Model* (LLM) that provides a stochastic residual correction to the  $k$ NN prior. Deterministic stabilisers ensure physical plausibility (e.g. monotonic material loss), while Monte Carlo rollouts yield pixel-wise uncertainty maps. The entire system is implemented as a modular Python package (`corrosion_forecast`) and evaluated with spatial overlap, boundary accuracy, and probabilistic calibration metrics.

## Contents

<b>I Mathematical Theory</b>	<b>2</b>
<b>1 Principal Component Analysis via Singular Value Decomposition</b>	<b>2</b>
1.1 Data Matrix Construction . . . . .	2
1.2 Mean Subtraction . . . . .	2
1.3 Singular Value Decomposition . . . . .	2
1.4 Truncated Reconstruction . . . . .	2
1.5 Explained Variance Ratio . . . . .	3
1.6 Randomised SVD . . . . .	3
1.7 Latent Normalisation . . . . .	3
<b>2 Signed Distance Field (SDF) Representation</b>	<b>3</b>
2.1 Definition . . . . .	3
2.2 Zero Level-Set Recovery . . . . .	4
2.3 Advantages over Raw Binary Masks . . . . .	4
<b>3 <math>k</math>-Nearest Neighbour Delta Prediction</b>	<b>4</b>
3.1 Feature Vector Construction . . . . .	4
3.2 Library Construction . . . . .	5
3.3 Inverse-Distance Weighted Prediction . . . . .	5
<b>4 LLM-in-the-Loop Forecasting</b>	<b>5</b>
4.1 Hybrid Architecture . . . . .	5
4.2 Prompt Structure . . . . .	5
4.3 Kinematic Context . . . . .	6

4.4	Rationale for the Hybrid Approach . . . . .	6
4.5	Robust JSON Parsing and Retry Logic . . . . .	6
<b>5</b>	<b>Deterministic Stabilisers</b>	<b>6</b>
5.1	Training-Derived Caps . . . . .	7
5.2	Horizon Damping . . . . .	7
5.3	Velocity-Relative Cap . . . . .	7
5.4	Monotonic SDF Shrinkage . . . . .	7
5.5	Optional SDF Smoothing . . . . .	7
5.6	Mask Post-Processing . . . . .	8
<b>6</b>	<b>Monte Carlo Uncertainty Quantification</b>	<b>8</b>
6.1	Stochastic Rollouts . . . . .	8
6.2	Probability Map . . . . .	8
6.3	Mean Prediction . . . . .	8
<b>7</b>	<b>Evaluation Metrics</b>	<b>8</b>
7.1	Spatial Overlap Metrics . . . . .	8
7.1.1	Intersection over Union (IoU) . . . . .	8
7.1.2	Sørensen–Dice Coefficient . . . . .	9
7.2	Boundary F1-Score (BF1) . . . . .	9
7.3	Probabilistic Calibration Metrics . . . . .	9
7.3.1	Brier Score . . . . .	9
7.3.2	Negative Log-Likelihood (NLL) . . . . .	9
7.3.3	Expected Calibration Error (ECE) . . . . .	9
<b>8</b>	<b>Auto-Tune PCA</b>	<b>10</b>
8.1	Motivation . . . . .	10
8.2	Grid Search Protocol . . . . .	10
8.3	Quality Thresholds . . . . .	10
8.4	Subspace Stability . . . . .	10
<b>II</b>	<b>Code Architecture and Documentation</b>	<b>10</b>
<b>9</b>	<b>Package Structure</b>	<b>10</b>
<b>10</b>	<b>Module Descriptions</b>	<b>11</b>
10.1	<code>config.py</code> — Global Configuration . . . . .	11
10.1.1	Key Constants . . . . .	11
10.1.2	Dependencies . . . . .	11
10.2	<code>data_loading.py</code> — Dataset I/O . . . . .	12
10.2.1	Key Functions . . . . .	12
10.2.2	Frame Caching . . . . .	12
10.2.3	Data Convention . . . . .	12
10.2.4	Dependencies . . . . .	12
10.3	<code>sdf_utils.py</code> — SDF Utilities . . . . .	12
10.3.1	Key Functions . . . . .	13
10.3.2	Post-Processing Cascade . . . . .	13
10.3.3	Dependencies . . . . .	13
10.4	<code>pca.py</code> — PCA Basis Fitting . . . . .	13
10.4.1	Key Functions . . . . .	13
10.4.2	Implementation Details . . . . .	14

10.4.3 Dependencies . . . . .	14
10.5 <code>metrics.py</code> — Evaluation Metrics . . . . .	14
10.5.1 Key Functions . . . . .	14
10.5.2 Dependencies . . . . .	15
10.6 <code>knn.py</code> — kNN Delta Predictor . . . . .	15
10.6.1 Key Functions . . . . .	15
10.6.2 Feature Construction . . . . .	15
10.6.3 Dependencies . . . . .	15
10.7 <code>llm_interface.py</code> — LLM API Interface . . . . .	15
10.7.1 Key Functions . . . . .	15
10.7.2 Retry and Recovery Strategy . . . . .	16
10.7.3 HTTP Details . . . . .	16
10.7.4 Dependencies . . . . .	17
10.8 <code>forecasting.py</code> — SDF Forecasting Engine . . . . .	17
10.8.1 Key Functions . . . . .	17
10.8.2 Single-Step Pipeline ( <code>forecast_next_sdf</code> ) . . . . .	18
10.8.3 Dependencies . . . . .	18
10.9 <code>autotune.py</code> — PCA Auto-Tuning . . . . .	18
10.9.1 Key Functions . . . . .	18
10.9.2 Grid Configuration . . . . .	19
10.9.3 Selection Logic . . . . .	19
10.9.4 Dependencies . . . . .	20
10.10 <code>ablation.py</code> — Ablation Study Runner . . . . .	20
10.10.1 Key Functions . . . . .	20
10.10.2 Ablation Variants . . . . .	20
10.10.3 Checkpoint/Resume . . . . .	20
10.10.4 Dependencies . . . . .	21
10.11 <code>plotting.py</code> — Visualisation Routines . . . . .	21
10.11.1 Available Plots . . . . .	21
10.11.2 Dependencies . . . . .	21
10.12 <code>main.py</code> — Pipeline Orchestrator . . . . .	21
10.12.1 Execution Sequence . . . . .	21
10.12.2 Usage . . . . .	22
10.12.3 Dependencies . . . . .	22
<b>11 Data Flow Diagram</b>	<b>22</b>
<b>12 Configuration and Reproducibility</b>	<b>23</b>
12.1 Environment Variables . . . . .	23
12.2 Modifying Hyperparameters . . . . .	23
12.3 Random Seeds . . . . .	23
12.4 Dataset Layout . . . . .	23
12.5 Reproducibility Checklist . . . . .	24

# Part I

## Mathematical Theory

### 1 Principal Component Analysis via Singular Value Decomposition

#### 1.1 Data Matrix Construction

Let  $T$  denote the total number of temporal snapshots and  $S$  the number of training slice indices. Each corrosion mask image of resolution  $H \times W$  is first converted to a field representation (Section 2) and then flattened into a row vector of length  $N = H \cdot W$ . Stacking all  $M = S \times T$  training fields yields the data matrix

$$\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^\top \\ \mathbf{x}_2^\top \\ \vdots \\ \mathbf{x}_M^\top \end{pmatrix} \in \mathbb{R}^{M \times N}. \quad (1)$$

#### 1.2 Mean Subtraction

The global mean field is

$$\boldsymbol{\mu} = \frac{1}{M} \sum_{i=1}^M \mathbf{x}_i \in \mathbb{R}^N, \quad (2)$$

and the mean-centred data matrix is

$$\mathbf{X}_c = \mathbf{X} - \mathbf{1}_M \boldsymbol{\mu}^\top \in \mathbb{R}^{M \times N}, \quad (3)$$

where  $\mathbf{1}_M$  is the  $M$ -vector of ones.

#### 1.3 Singular Value Decomposition

The (economy) SVD of the centred matrix is

$$\mathbf{X}_c = \mathbf{U} \mathbf{S} \mathbf{V}^\top, \quad (4)$$

where  $\mathbf{U} \in \mathbb{R}^{M \times M}$  contains left singular vectors,  $\mathbf{S} = \text{diag}(\sigma_1, \dots, \sigma_{\min(M,N)})$  are singular values in descending order, and  $\mathbf{V} \in \mathbb{R}^{N \times N}$  contains right singular vectors. The columns of  $\mathbf{V}$  (equivalently the rows of  $\mathbf{V}^\top$ ) are the *principal component directions*  $\mathbf{v}_1, \dots, \mathbf{v}_N$ .

#### 1.4 Truncated Reconstruction

Retaining only  $K \ll \min(M, N)$  components, an arbitrary field  $\mathbf{x}$  is reconstructed as

$$\hat{\mathbf{x}} = \boldsymbol{\mu} + \sum_{k=1}^K z_k \mathbf{v}_k, \quad (5)$$

where the *latent coefficients* are

$$z_k = (\mathbf{x} - \boldsymbol{\mu})^\top \mathbf{v}_k, \quad k = 1, \dots, K. \quad (6)$$

The latent vector  $\mathbf{z} = (z_1, \dots, z_K)^\top \in \mathbb{R}^K$  provides a low-dimensional representation of the corrosion field.

## 1.5 Explained Variance Ratio

Each singular value  $\sigma_k$  is related to the eigenvalue of the covariance matrix via  $\lambda_k = \sigma_k^2/(M - 1)$ . The fraction of variance explained by the first  $K$  components is

$$\text{EVR}(K) = \frac{\sum_{k=1}^K \lambda_k}{\min(M,N)} = \frac{\sum_{k=1}^K \sigma_k^2}{\min(M,N)}. \quad (7)$$

$$\sum_{j=1}^K \lambda_j \quad \sum_{j=1}^K \sigma_j^2$$

## 1.6 Randomised SVD

For large  $M$  and  $N$ , the full SVD is computationally prohibitive. Following [1], we use the *randomised SVD* algorithm, which approximates the top- $K$  singular triplets in  $\mathcal{O}(MN \log K)$  time rather than  $\mathcal{O}(MN \min(M, N))$ . The algorithm proceeds by:

1. Drawing a random Gaussian matrix  $\Omega \in \mathbb{R}^{N \times (K+p)}$ , where  $p$  is a small over-sampling parameter.
2. Computing the sample matrix  $\mathbf{Y} = \mathbf{X}_c \Omega$ .
3. Obtaining an orthonormal basis  $\mathbf{Q}$  for the column space of  $\mathbf{Y}$  via QR decomposition.
4. Forming the small matrix  $\mathbf{B} = \mathbf{Q}^\top \mathbf{X}_c$  and computing its SVD.

In our implementation, we use `sklearn.utils.extmath.randomized_svd` when available, falling back to `numpy.linalg.svd` otherwise.

## 1.7 Latent Normalisation

To ensure numerical stability when feeding latent vectors to downstream predictors, we compute per-component statistics

$$\bar{z}_k = \frac{1}{M} \sum_{i=1}^M z_k^{(i)}, \quad s_k = \sqrt{\frac{1}{M} \sum_{i=1}^M (z_k^{(i)} - \bar{z}_k)^2 + \epsilon}, \quad (8)$$

with  $\epsilon = 10^{-6}$ , and work with the normalised latent  $\tilde{z}_k = (z_k - \bar{z}_k)/s_k$  throughout the forecasting pipeline.

## 2 Signed Distance Field (SDF) Representation

### 2.1 Definition

**Definition 2.1** (Signed Distance Field). Given a binary material mask  $\mathcal{M} \subseteq \mathbb{R}^2$  (the set of pixels labelled as material), the Signed Distance Field is defined at every pixel location  $\mathbf{x}$  as

$$\text{SDF}(\mathbf{x}) = d_{\text{outside}}(\mathbf{x}) - d_{\text{inside}}(\mathbf{x}), \quad (9)$$

where

$$d_{\text{inside}}(\mathbf{x}) = \inf_{\mathbf{x} \in \partial\mathcal{M}} \|\mathbf{x} - \mathbf{x}\|_2 \quad \text{for } \mathbf{x} \in \mathcal{M}, \quad (10)$$

$$d_{\text{outside}}(\mathbf{x}) = \inf_{\mathbf{x} \in \partial\mathcal{M}} \|\mathbf{x} - \mathbf{x}\|_2 \quad \text{for } \mathbf{x} \notin \mathcal{M}, \quad (11)$$

and  $\partial\mathcal{M}$  denotes the material boundary.

By this convention:

- $\text{SDF}(\mathbf{x}) < 0$  for pixels **inside** the material.
- $\text{SDF}(\mathbf{x}) > 0$  for pixels **outside** the material.
- $\text{SDF}(\mathbf{x}) = 0$  at the material boundary.

In practice, both distance transforms are computed using the Euclidean Distance Transform (`scipy.ndimage.distance_transform_edt`).

## 2.2 Zero Level-Set Recovery

The material mask can always be recovered from the SDF by thresholding:

$$\mathcal{M} = \{\mathbf{x} : \text{SDF}(\mathbf{x}) \leq 0\}. \quad (12)$$

## 2.3 Advantages over Raw Binary Masks

Using SDF representations instead of raw  $\{0, 255\}$  binary masks for PCA confers several advantages:

1. **Smoothness:** The SDF is a continuous, piecewise-smooth function. Small perturbations of the boundary produce small perturbations of the SDF in the  $L^2$  norm, whereas binary masks exhibit discontinuous jumps.
2. **Boundary sensitivity:** The SDF gradient is largest near the boundary (magnitude = 1 everywhere for exact distance functions), ensuring that PCA modes concentrate representational power on the shape of the boundary rather than on interior fill.
3. **Linear interpolation:** Averaging two SDF fields in PCA latent space yields a geometrically meaningful intermediate shape, whereas averaging two binary masks produces ambiguous grey values.
4. **Reconstruction quality:** Truncated PCA reconstruction of a smooth SDF produces a smooth field whose zero level-set gives a clean boundary; reconstructing a binary mask introduces ringing artefacts.

## 3 $k$ -Nearest Neighbour Delta Prediction

### 3.1 Feature Vector Construction

The kNN predictor operates in the normalised PCA latent space. At time  $t$ , the feature vector for a query is composed of:

$$\mathbf{f} = [\tilde{\mathbf{z}}_t, \dot{\tilde{\mathbf{z}}}_t, t_{\text{norm}}, \Delta t_{\text{norm}}] \in \mathbb{R}^{2K+2}, \quad (13)$$

where:

- $\tilde{\mathbf{z}}_t \in \mathbb{R}^K$  is the normalised latent state at time  $t$ .
- $\dot{\tilde{\mathbf{z}}}_t = \tilde{\mathbf{z}}_t - \tilde{\mathbf{z}}_{t-1}$  is the latent velocity (zero for  $t = 0$ ).
- $t_{\text{norm}} = t_h/t_{\max}$  and  $\Delta t_{\text{norm}} = \Delta t_h/t_{\max}$  are the normalised current time and step size (in hours), respectively.

### 3.2 Library Construction

For each training slice and each pair of consecutive time-steps  $(t, t + 1)$ , we store the pair  $(\mathbf{f}_t, \boldsymbol{\delta}_t)$  where  $\boldsymbol{\delta}_t = \tilde{\mathbf{z}}_{t+1} - \tilde{\mathbf{z}}_t$  is the ground-truth latent delta. The full library is

$$\mathcal{L} = \{(\mathbf{f}^{(i)}, \boldsymbol{\delta}^{(i)})\}_{i=1}^{|\mathcal{L}|}. \quad (14)$$

### 3.3 Inverse-Distance Weighted Prediction

Given a query feature  $\mathbf{f}_q$ , we find the  $k$  nearest neighbours  $\mathcal{N}_k(\mathbf{f}_q)$  in  $\mathcal{L}$  by Euclidean distance. For each neighbour  $i \in \mathcal{N}_k$ , the distance is

$$d_i = \|\mathbf{f}_q - \mathbf{f}^{(i)}\|_2. \quad (15)$$

The inverse-distance weights are

$$w_i = \frac{1/d_i}{\sum_{j \in \mathcal{N}_k} 1/d_j}, \quad i \in \mathcal{N}_k, \quad (16)$$

with a small additive constant  $\epsilon = 10^{-6}$  for numerical stability.

The **weighted mean** (the kNN prior) is

$$\boldsymbol{\delta}_{\text{prior}} = \boldsymbol{\mu}_{\text{kNN}} = \sum_{i \in \mathcal{N}_k} w_i \boldsymbol{\delta}^{(i)}, \quad (17)$$

and the **weighted variance** (for uncertainty estimation) is

$$\sigma_{\text{kNN}}^2 = \sum_{i \in \mathcal{N}_k} w_i (\boldsymbol{\delta}^{(i)} - \boldsymbol{\mu}_{\text{kNN}})^2. \quad (18)$$

## 4 LLM-in-the-Loop Forecasting

### 4.1 Hybrid Architecture

The pipeline employs a *hybrid* forecasting strategy: the kNN module provides a physics-informed, data-driven baseline prediction (the *prior*), and a Large Language Model (LLM) adds a stochastic *residual correction* that captures non-linear dependencies beyond nearest-neighbour interpolation.

The final predicted delta is

$$\boldsymbol{\delta} = \boldsymbol{\delta}_{\text{prior}} + \alpha \boldsymbol{\delta}_{\text{LLM}}, \quad (19)$$

where  $\alpha \in [0, 1]$  is the `RESIDUAL_SCALE` parameter (default  $\alpha = 0.7$ ) and  $\boldsymbol{\delta}_{\text{LLM}}$  is the residual vector returned by the LLM.

### 4.2 Prompt Structure

The LLM receives a single **JSON payload** as the user message, containing:

```

1 {
2     "D": <int: dimension of latent space>,
3     "dt_hours": <float: time-step in hours>,
4     "last_latent": [z_1, z_2, ..., z_K],
5     "velocity": [v_1, v_2, ..., v_K],
6     "acceleration": [a_1, a_2, ..., a_K],
7     "delta_prior": [d_1, d_2, ..., d_K],
8     "prior_std": [s_1, s_2, ..., s_K],
9     "residual_norm_cap": <float>,
10    "residual_comp_cap": <float>,
11    "rollout_nonce": <float: for stochastic diversity>,
12    "metadata_context": { ... }
13 }
```

The system prompt instructs the model to return **only** valid JSON with a single key `predicted_delta` containing exactly  $K$  floats. Explicitly:

```
Return ONLY JSON. No prose. No markdown.
Output ONLY JSON with exactly one key predicted_delta.
predicted_delta must be a list of exactly D=<K> floats.
Interpret predicted_delta as a RESIDUAL to add to delta_prior.
Keep it small: L2(residual) <= residual_norm_cap
    and abs(component) <= residual_comp_cap.
```

### 4.3 Kinematic Context

To help the LLM reason about the dynamics, we provide:

- The **last latent state**  $\tilde{\mathbf{z}}_t$ .
- The **velocity**  $\dot{\tilde{\mathbf{z}}}_t = \tilde{\mathbf{z}}_t - \tilde{\mathbf{z}}_{t-1}$ .
- The **acceleration**  $\ddot{\tilde{\mathbf{z}}}_t = \dot{\tilde{\mathbf{z}}}_t - \dot{\tilde{\mathbf{z}}}_{t-1}$  (zero if fewer than three observations are available).

### 4.4 Rationale for the Hybrid Approach

1. **kNN provides structure:** The inverse-distance-weighted neighbourhood average is fast, purely data-driven, and guaranteed to stay within the convex hull of training deltas. It encodes the *typical* magnitude and direction of latent evolution.
2. **LLM adds flexibility:** By operating as a residual predictor, the LLM need only learn the *correction* to the prior, which is typically much smaller than the full delta. This reduces the effective prediction complexity and makes the system more robust to LLM hallucinations.
3. **Bounded risk:** The residual is capped both in norm and per component (see Section 5), so even a poorly calibrated LLM output cannot drive the prediction far from the kNN baseline.

### 4.5 Robust JSON Parsing and Retry Logic

LLM outputs may deviate from the expected format. The parsing module implements a multi-stage recovery strategy:

1. **Strict parsing:** attempt `json.loads` on the raw response.
2. **Regex extraction:** search for an embedded `{...}` JSON blob.
3. **Text salvage:** if enabled, extract floating-point numbers from potentially truncated text after the key `predicted_delta`.
4. **Length repair:** trim or zero-pad the vector if its length does not match  $K$ .
5. **Conversational retry:** append the malformed response as an assistant turn and ask the model to correct itself (up to `LLM_MAX_RETRIES` times).

## 5 Deterministic Stabilisers

After obtaining the combined delta  $\delta$  (Equation 19), several deterministic constraints are applied to ensure physical plausibility and prevent forecast divergence.

## 5.1 Training-Derived Caps

From all consecutive-step deltas in the training set, we compute:

- **Magnitude cap:** the 95th percentile of  $\|\delta^{(i)}\|_2$  across training pairs, scaled by a multiplier  $\rho_{\text{mag}}$  (default 1.2):

$$c_{\text{mag}} = \rho_{\text{mag}} \cdot P_{95}(\{\|\delta^{(i)}\|_2\}). \quad (20)$$

If  $\|\delta\|_2 > c_{\text{mag}}$ , the delta is rescaled:  $\delta \leftarrow \delta \cdot c_{\text{mag}} / \|\delta\|_2$ .

- **Per-component cap:** the 99th percentile of  $|\delta_k^{(i)}|$  for each component  $k$ , scaled by  $\rho_{\text{comp}}$  (default 1.2):

$$c_k = \rho_{\text{comp}} \cdot P_{99}(|\delta_k^{(i)}|). \quad (21)$$

Each component is clipped:  $\delta_k \leftarrow \text{clip}(\delta_k, -c_k, c_k)$ .

## 5.2 Horizon Damping

For multi-step autoregressive rollouts, prediction uncertainty compounds with each step. To counteract drift, the delta magnitude is attenuated exponentially with the forecast horizon  $h$ :

$$\delta \leftarrow \delta \cdot \gamma^{(h-1)}, \quad (22)$$

where  $\gamma \in (0, 1]$  is the damping factor (default  $\gamma = 0.95$ ). The first horizon step ( $h = 1$ ) is unmodified.

## 5.3 Velocity-Relative Cap

To prevent the predicted change from being unreasonably large compared to the observed rate of evolution, we impose

$$\|\delta\|_2 \leq \alpha_v \|\dot{\mathbf{z}}\|_2 + \beta_v, \quad (23)$$

where  $\alpha_v$  (default 6.0) and  $\beta_v$  (default 0.2) are tuneable constants. If the constraint is violated,  $\delta$  is rescaled to satisfy it. A minimum cap floor of  $0.25 \cdot c_{\text{mag}}$  prevents the constraint from being too restrictive when velocity is very small.

**Remark 5.1.** When the kNN guide is active, the velocity-relative cap is optionally disabled (`DISABLE_VEL_CAP_WHEN_KNN=True`) because the kNN prior already implicitly regularises the delta magnitude.

## 5.4 Monotonic SDF Shrinkage

Corrosion is an irreversible process: material can only be removed, never added. In the SDF representation, material loss corresponds to the SDF becoming more positive (or less negative). We enforce this via a pixel-wise maximum:

$$\text{SDF}_{t+1}(\mathbf{x}) = \max(\text{SDF}_{\text{pred}}(\mathbf{x}), \text{SDF}_t(\mathbf{x})) \quad \forall \mathbf{x}. \quad (24)$$

Since  $\text{SDF}(\mathbf{x}) \leq 0$  indicates material and the maximum operator only *increases* SDF values, this guarantees  $\mathcal{M}_{t+1} \subseteq \mathcal{M}_t$  (monotonic shrinkage).

## 5.5 Optional SDF Smoothing

Before thresholding, the predicted SDF may be lightly smoothed with a Gaussian filter ( $\sigma = 0.8$  px) to suppress high-frequency artefacts introduced by PCA truncation:

$$\text{SDF}_{\text{smooth}} = G_\sigma * \text{SDF}_{\text{pred}}, \quad (25)$$

where  $G_\sigma$  is the Gaussian kernel.

## 5.6 Mask Post-Processing

After thresholding the SDF to obtain a binary mask, morphological post-processing is applied in the following order:

1. **Fill holes:** close interior voids using `binary_fill_holes`.
2. **Remove small components:** discard connected components with fewer than 200 pixels.
3. **Largest component:** keep only the largest connected component (the alloy wire cross-section).
4. **Monotonic mask shrinkage:** intersect the predicted mask with the previous-step mask:  $\mathcal{M}_{t+1} \leftarrow \mathcal{M}_{t+1} \cap \mathcal{M}_t$ .

## 6 Monte Carlo Uncertainty Quantification

### 6.1 Stochastic Rollouts

With the LLM temperature  $\tau > 0$ , each call to the language model produces a different residual correction, inducing stochasticity in the forecast. We perform  $R$  independent autoregressive rollouts (default  $R = 8$ ), each seeded with a distinct `rollout_nonce`. The LLM effective temperature decays across horizons:

$$\tau_{\text{eff}}(h) = \tau_0 \cdot \eta^{\max(0, h-1)}, \quad (26)$$

where  $\eta$  (default 0.98) is the horizon temperature decay.

### 6.2 Probability Map

The pixel-wise probability of material survival is estimated from the ensemble of rollout predictions:

$$P(\text{material} | \mathbf{x}) = \frac{1}{R} \sum_{r=1}^R \mathbb{I}[\mathbf{x} \in \mathcal{M}^{(r)}], \quad (27)$$

where  $\mathcal{M}^{(r)}$  is the predicted material mask from rollout  $r$  and  $\mathbb{I}[\cdot]$  is the indicator function.

### 6.3 Mean Prediction

The consensus (mean) prediction thresholds the probability map at 0.5:

$$\hat{\mathcal{M}} = \{\mathbf{x} : P(\text{material} | \mathbf{x}) \geq 0.5\}. \quad (28)$$

**Remark 6.1.** Regions where  $P(\text{material} | \mathbf{x})$  is close to 0.5 indicate high uncertainty. These typically localise along the advancing corrosion front, providing physically meaningful uncertainty estimates.

## 7 Evaluation Metrics

### 7.1 Spatial Overlap Metrics

Let  $A$  and  $B$  denote the ground-truth and predicted material masks, respectively.

#### 7.1.1 Intersection over Union (IoU)

$$\text{IoU}(A, B) = \frac{|A \cap B|}{|A \cup B|}, \quad (29)$$

where  $|\cdot|$  denotes the pixel count.

### 7.1.2 Sørensen–Dice Coefficient

$$\text{Dice}(A, B) = \frac{2|A \cap B|}{|A| + |B|}. \quad (30)$$

**Remark 7.1.** IoU and Dice are related by  $\text{Dice} = 2\text{IoU}/(1 + \text{IoU})$ , so Dice is always  $\geq$  IoU for the same prediction. IoU is the stricter metric.

## 7.2 Boundary F1-Score (BF1)

The boundary F1-score evaluates boundary localisation accuracy. First, extract boundary pixels of each mask via morphological erosion:

$$\partial A = A \oplus (A \ominus B), \quad (31)$$

where  $\ominus$  is erosion and  $\oplus$  is symmetric difference (XOR). With a tolerance of  $\tau$  pixels (implemented as dilation by  $\tau$  iterations):

$$\text{Prec}_\tau = \frac{|\partial B \cap D_\tau(\partial A)|}{|\partial B|}, \quad (32)$$

$$\text{Rec}_\tau = \frac{|\partial A \cap D_\tau(\partial B)|}{|\partial A|}, \quad (33)$$

where  $D_\tau(\cdot)$  denotes dilation by  $\tau$  pixels. The Boundary F1-score is

$$\text{BF1}_\tau = \frac{2\text{Prec}_\tau \text{Rec}_\tau}{\text{Prec}_\tau + \text{Rec}_\tau}. \quad (34)$$

## 7.3 Probabilistic Calibration Metrics

Let  $p_i \in [0, 1]$  be the predicted probability that pixel  $i$  is material, and  $y_i \in \{0, 1\}$  the ground-truth label.

### 7.3.1 Brier Score

$$\text{BS} = \frac{1}{N} \sum_{i=1}^N (p_i - y_i)^2. \quad (35)$$

Lower is better; a perfect deterministic predictor achieves  $\text{BS} = 0$ .

### 7.3.2 Negative Log-Likelihood (NLL)

$$\text{NLL} = -\frac{1}{N} \sum_{i=1}^N [y_i \log(p_i + \epsilon) + (1 - y_i) \log(1 - p_i + \epsilon)], \quad (36)$$

with  $\epsilon = 10^{-6}$  for numerical stability.

### 7.3.3 Expected Calibration Error (ECE)

Partition pixels into  $B$  bins by predicted probability. Let  $\mathcal{B}_b$  be the set of pixels in bin  $b$ , with  $n_b = |\mathcal{B}_b|$ , mean confidence  $\bar{p}_b$ , and empirical accuracy  $\bar{y}_b$ . Then:

$$\text{ECE} = \sum_{b=1}^B \frac{n_b}{N} |\bar{y}_b - \bar{p}_b|. \quad (37)$$

ECE is visualised using a *reliability diagram* (Section 10.11).

## 8 Auto-Tune PCA

### 8.1 Motivation

The number of principal components  $K$  and the training set size jointly determine reconstruction fidelity. Too few components lose fine-grained boundary detail; too many overfit noise and inflate the latent dimension for downstream predictors.

### 8.2 Grid Search Protocol

The auto-tune procedure performs a grid search over  $(\text{n\_train}, K)$  using *oracle reconstruction* on a held-out validation set. Oracle reconstruction projects the ground-truth field onto the PCA basis with  $K$  components and then thresholds to obtain a material mask:

$$\hat{\mathbf{x}}^{(K)} = \boldsymbol{\mu} + \sum_{k=1}^K [(\mathbf{x} - \boldsymbol{\mu})^\top \mathbf{v}_k] \mathbf{v}_k, \quad \hat{\mathcal{M}}^{(K)} = \{\mathbf{x} : \hat{\mathbf{x}}^{(K)}(\mathbf{x}) \leq 0\}. \quad (38)$$

To reduce computational cost, all fields are spatially downsampled by a factor of 2 (default), and only a subset of time-steps  $\{0, 3, 6, 9\}$  is used.

### 8.3 Quality Thresholds

The smallest  $K$  satisfying *all* of the following is selected as  $K_{\text{LLM}}$ :

$$\overline{\text{IoU}}(K) \geq 0.985, \quad (39)$$

$$\overline{\text{Dice}}(K) \geq 0.990, \quad (40)$$

$$\overline{\text{BF1}}(K) \geq 0.930. \quad (41)$$

The reconstruction dimensionality  $K_{\text{recon}} \geq K_{\text{LLM}}$  is selected as the point of diminishing returns, where the marginal IoU and Dice gain from an additional component drops below  $5 \times 10^{-4}$ .

### 8.4 Subspace Stability

To verify that the PCA basis is robust to training set composition, we fit two bases on independent random subsets and measure the *principal angles* between the resulting  $K$ -dimensional subspaces:

$$\theta_j = \arccos(\sigma_j(\mathbf{V}_A^{(K)\top} \mathbf{V}_B^{(K)})), \quad j = 1, \dots, K, \quad (42)$$

where  $\sigma_j(\cdot)$  are singular values. A candidate is accepted only if  $\bar{\theta} \leq 8^\circ$  (mean principal angle).

## Part II

# Code Architecture and Documentation

## 9 Package Structure

The pipeline is implemented as a single Python package with the following layout:

```
corrosion_forecast/
|-- __init__.py
|-- config.py
|-- data_loading.py
|-- sdf_utils.py
|-- pca.py
|-- metrics.py
```

```

|-- knn.py
|-- llm_interface.py
|-- forecasting.py
|-- autotune.py
|-- ablation.py
|-- plotting.py
|-- main.py
`-- requirements.txt

```

Dependencies:

```

numpy >=1.24      scipy >=1.10      matplotlib >=3.7
pandas >=2.0       imageio >=2.31     requests >=2.28
scikit-learn >=1.3

```

## 10 Module Descriptions

### 10.1 config.py — Global Configuration

**Purpose:** Centralises all tuneable hyperparameters, file-system paths, and feature flags so that experiments are reproducible and modifications require editing only a single file.

#### 10.1.1 Key Constants

Group	Parameter	Default
Paths	BASE_DIR	./Corrosion_Masks (or \$CORROSION_BASE_DIR)
	METADATA_CSV	<BASE_DIR>/metadata.csv
	NUM_TIMESTEPS	10
LLM	OPENROUTER_API_KEY	\$OPENROUTER_API_KEY
	LLM_MODEL	openai/gpt-4o-mini
	LLM_TEMPERATURE	0.6
	LLM_MAX_TOKENS	600
	LLM_MAX_RETRIES	2
PCA	K_LLM	40
	K_PCS_MAX	80
	NUM_TRAIN_SLICES	400
kNN	KNN_K	16
	KNN_MAX_SLICES	300
	RESIDUAL_SCALE	0.7
Stabilisers	HORIZON_DAMP	0.95
	VEL_REL_MULT	6.0
	VEL_REL_BIAS	0.20
Experiment	NUM_TEST_SLICES	2
	TEST_START_TIMES	[3, 5]
	NUM_ROLLOUTS	8

Table 1: Selected configuration parameters and their defaults.

#### 10.1.2 Dependencies

Standard library only (`os`), plus `numpy` for dtype constants.

## 10.2 data\_loading.py — Dataset I/O

**Purpose:** Handles all file-system interaction: reading TIFF corrosion mask slices, parsing the metadata CSV, binarisation, and in-memory frame caching.

### 10.2.1 Key Functions

```
1 def load_frame_uint8(
2     slice_idx: int, t: int, use_cache: bool = True
3 ) -> Optional[np.ndarray]:
4     """Load a single grayscale frame as uint8.
5     Returns None if the file does not exist."""
6
7 def load_slice_across_time(
8     slice_idx: int, use_cache: bool = False
9 ) -> Optional[List[np.ndarray]]:
10    """Load one slice index across all NUM_TIMESTEPS time-steps.
11    Returns list of (H, W) uint8 images, or None if any missing."""
12
13 def get_total_slice_count() -> int:
14    """Count .tif files in the first time-step folder."""
15
16 def load_times_from_metadata() -> np.ndarray:
17    """Return degradation times (hours) from metadata.csv."""
18
19 def binarize_0_255(img: np.ndarray, thr: int = 128) -> np.ndarray:
20    """Threshold to {0, 255}."""
21
22 def to_material_bool(img_0_255: np.ndarray) -> np.ndarray:
23    """Convert {0,255} mask to boolean (True = material)."""
24
25 def from_material_bool(material_bool: np.ndarray) -> np.ndarray:
26    """Inverse of to_material_bool."""
27
28 def measure_area_material(material_bool: np.ndarray) -> int:
29    """Count material pixels."""
```

### 10.2.2 Frame Caching

A module-level dictionary `_FRAME_CACHE` maps `(slice_idx, t)` to the loaded `uint8` array. The cache is transparent to callers and can be cleared via `clear_frame_cache()`.

### 10.2.3 Data Convention

In the TIFF files, pixel value 0 (black) denotes **material** and pixel value 255 (white) denotes **background/void**. The helper `to_material_bool` implements `img == 0`.

### 10.2.4 Dependencies

`imageio, numpy, pandas, config.`

## 10.3 sdf\_utils.py — SDF Utilities

**Purpose:** Converts between boolean material masks and Signed Distance Fields, and provides morphological post-processing routines.

### 10.3.1 Key Functions

```
1 def material_to_sdf(material_bool: np.ndarray) -> np.ndarray:
2     """Boolean mask -> SDF (float32). Negative inside material."""
3
4 def sdf_to_material(sdf: np.ndarray) -> np.ndarray:
5     """SDF -> boolean mask via zero level-set (SDF <= 0)."""
6
7 def keep_largest_component(mask_bool: np.ndarray) -> np.ndarray:
8     """Keep only the largest connected component."""
9
10 def remove_small_components(
11     mask_bool: np.ndarray, min_pixels: int = 200
12 ) -> np.ndarray:
13     """Remove components smaller than min_pixels."""
14
15 def postprocess_material(
16     pred_material: np.ndarray,
17     prev_material: np.ndarray = None,
18 ) -> np.ndarray:
19     """Apply cascade: fill holes -> remove small ->
20     keep largest -> monotonic shrink."""
21
22 def smooth_sdf(sdf: np.ndarray) -> np.ndarray:
23     """Optional Gaussian smoothing (sigma from config)."""
```

### 10.3.2 Post-Processing Cascade

The `postprocess_material` function applies, in order:

1. `binary_fill_holes` (if `FILL_HOLES`).
2. `remove_small_components` (if `MIN_COMPONENT_PIXELS > 0`).
3. `keep_largest_component` (if `ENFORCE_SINGLE_COMPONENT`).
4. Monotonic shrinkage via boolean AND with `prev_material` (if `ENFORCE_MONOTONIC_SHRINK`).

Each step is gated by the corresponding configuration flag.

### 10.3.3 Dependencies

`scipy.ndimage`, `numpy`, `config`.

## 10.4 pca.py — PCA Basis Fitting

**Purpose:** Fits a global PCA basis from training slices via (randomised) SVD, and provides projection/reconstruction utilities.

### 10.4.1 Key Functions

```
1 def fit_global_pca(
2     num_slices: int,
3     max_index: int,
4     k_max: int = 60,
5     seed: int = 0,
6 ) -> Tuple[np.ndarray, np.ndarray, np.ndarray, np.ndarray]:
7     """Fit PCA basis from training slices.
8     Returns: (mean_field, Vt, z_mean, z_std)."""
9
```

```

10 def project_field(
11     field_2d: np.ndarray, mean_field: np.ndarray, Vt: np.ndarray
12 ) -> np.ndarray:
13     """Project a 2D field onto PCA basis -> latent z."""
14
15 def reconstruct_field(
16     z: np.ndarray, mean_field: np.ndarray,
17     Vt: np.ndarray, H: int, W: int
18 ) -> np.ndarray:
19     """Reconstruct a 2D field from latent vector z."""

```

#### 10.4.2 Implementation Details

- Training slice IDs are randomly sampled up to `num_slices` from IDs in  $[0, \text{max\_index}]$ .
- For each slice, *all*  $T$  time-steps are included in the training matrix, yielding up to  $T \times \text{num\_slices}$  rows.
- If `scikit-learn` is available, `randomized_svd` is used; otherwise the code falls back to full SVD.
- The function returns both the PCA basis (`mean_field`, `Vt`) and the latent normalisation statistics (`z_mean`, `z_std`).

#### 10.4.3 Dependencies

`numpy`, `sklearn` (optional), `data_loading`, `sdf_utils`, `config`.

### 10.5 metrics.py — Evaluation Metrics

**Purpose:** Implements all evaluation metrics used for mask quality assessment and probabilistic calibration.

#### 10.5.1 Key Functions

```

1 def iou(gt: np.ndarray, pr: np.ndarray) -> float:
2     """Intersection over Union for boolean masks."""
3
4 def dice(gt: np.ndarray, pr: np.ndarray) -> float:
5     """Sorenson-Dice coefficient for boolean masks."""
6
7 def boundary_f1(
8     gt: np.ndarray, pr: np.ndarray, tol: int = 1
9 ) -> float:
10    """Boundary F1-score with tolerance (dilation)."""
11
12 def brier_score(probs: np.ndarray, y: np.ndarray) -> float:
13    """Mean squared error between probabilities and labels."""
14
15 def binary_nll(
16     probs: np.ndarray, y: np.ndarray, eps: float = 1e-6
17 ) -> float:
18    """Negative log-likelihood for binary predictions."""
19
20 def expected_calibration_error(
21     probs: np.ndarray, y: np.ndarray, n_bins: int = 15
22 ) -> float:
23    """Expected Calibration Error (ECE)."""

```

### 10.5.2 Dependencies

`numpy, scipy.ndimage` (for boundary extraction in BF1).

## 10.6 knn.py — kNN Delta Predictor

**Purpose:** Constructs a reference library of (feature, delta) pairs from training sequences and provides inverse-distance-weighted  $k$ -nearest neighbour predictions in PCA latent space.

### 10.6.1 Key Functions

```
1 def build_knn_library(
2     train_ids: list,
3     mean_field: np.ndarray,
4     Vt: np.ndarray,
5     z_mean: np.ndarray,
6     z_std: np.ndarray,
7     k_llm: int,
8     times_h_all: np.ndarray,
9     max_slices: int = 300,
10    seed: int = 0,
11 ) -> Dict[str, np.ndarray]:
12     """Build kNN library. Returns dict with keys
13     'X' (features), 'Y' (deltas), 'tmax'."""
14
15 def knn_predict_delta(
16     knn_lib: Dict[str, np.ndarray],
17     z_last: np.ndarray,
18     vel: np.ndarray,
19     t_hours: float,
20     dt_hours: float,
21     k: int = 16,
22 ) -> Tuple[np.ndarray, np.ndarray]:
23     """Predict next-step delta via IDW kNN.
24     Returns (mean_delta, std_delta)."""
25
```

### 10.6.2 Feature Construction

The internal function `_feat_from_state` builds the feature vector as defined in Equation (13): the last  $K$ -dimensional normalised latent, the velocity vector, and two normalised time scalars. This yields a feature of dimension  $2K + 2$ .

### 10.6.3 Dependencies

`numpy, data_loading, pca, config.`

## 10.7 llm\_interface.py — LLM API Interface

**Purpose:** Handles prompt construction, HTTP communication with the OpenRouter API, and robust JSON parsing of LLM responses.

### 10.7.1 Key Functions

```
1 def build_llm_prompt(
2     latents_norm_hist_top: np.ndarray,
3     dt_h: float,
4     rollout_nonce: Optional[float] = None,
5     meta_row: Optional[Dict] = None,
```

```

6     delta_prior: Optional[np.ndarray] = None,
7     prior_std: Optional[np.ndarray] = None,
8     residual_norm_cap: Optional[float] = None,
9     residual_comp_cap: Optional[float] = None,
10 ) -> str:
11     """Construct the user-message prompt as JSON payload."""
12
13 def parse_llm_delta(
14     text: str,
15     D: int,
16     allow_length_repair: bool = False,
17     allow_text_salvage: bool = False,
18 ) -> Tuple[Optional[np.ndarray], Dict]:
19     """Parse LLM response -> delta vector of length D.
20     Multi-stage: strict JSON, regex, salvage, length repair."""
21
22 def call_llm_delta(
23     latents_norm_hist_top: np.ndarray,
24     dt_h: float,
25     rollout_nonce: Optional[float] = None,
26     meta_row: Optional[Dict] = None,
27     horizon: int = 1,
28     delta_prior: Optional[np.ndarray] = None,
29     prior_std: Optional[np.ndarray] = None,
30     residual_norm_cap: Optional[float] = None,
31     residual_comp_cap: Optional[float] = None,
32 ) -> Tuple[np.ndarray, str, int, Dict]:
33     """Full LLM call with retries.
34     Returns (delta_vec, raw_text, http_status, info)."""

```

### 10.7.2 Retry and Recovery Strategy

The call loop (up to `LLM_MAX_RETRIES+1` attempts) follows Algorithm 1.

---

#### Algorithm 1 LLM call with progressive recovery

---

```

1: for attempt = 0, 1, ..., MAX_ATTEMPTS do
2:   Send prompt to LLM via HTTP POST
3:   Parse response with strict JSON
4:   if parse succeeds and length matches K then
5:     return parsed delta
6:   end if
7:   if length mismatch and retries remain then
8:     Append repair message; continue
9:   end if
10:  if parse failure and retries remain then
11:    Append correction message; continue
12:  end if
13:  Attempt text salvage and length repair (last resort)
14:  if salvage succeeds then return salvaged delta
15:  end if
16:  raise RuntimeError
17: end for

```

---

### 10.7.3 HTTP Details

- Endpoint: <https://openrouter.ai/api/v1/chat/completions>.

- Authentication: Bearer token via OPENROUTER\_API\_KEY.
- The `response_format` is set to `{"type": "json_object"}` when `USE_RESPONSE_FORMAT_JSON=True`, requesting structured output.
- The effective temperature decays with horizon:  $\tau_{\text{eff}} = \tau_0 \cdot 0.98^{h-1}$ .

#### 10.7.4 Dependencies

`json, re, requests, numpy, config.`

### 10.8 forecasting.py — SDF Forecasting Engine

**Purpose:** Implements the single-step SDF forecast (kNN prior + LLM residual + stabilisers) and the Monte Carlo rollout loop for multi-step autoregressive prediction.

#### 10.8.1 Key Functions

```

1 def compute_training_caps(
2     train_ids: List[int],
3     mean_field: np.ndarray,
4     Vt: np.ndarray,
5     z_mean: np.ndarray,
6     z_std: np.ndarray,
7     k_llm: int,
8     max_slices: int = 250,
9     seed: int = 0,
10 ) -> Dict[str, np.ndarray]:
11     """Compute magnitude and per-component caps from
12     training GT deltas."""
13
14 def clip_delta_to_training(
15     delta: np.ndarray,
16     caps: Optional[Dict],
17     horizon: int = 1,
18 ) -> np.ndarray:
19     """Apply training caps + horizon damping."""
20
21 def apply_velocity_relative_cap(
22     delta: np.ndarray,
23     vel: np.ndarray,
24     mult: float = 6.0,
25     bias: float = 0.2,
26     min_cap: Optional[float] = None,
27 ) -> np.ndarray:
28     """Cap delta magnitude relative to velocity."""
29
30 def forecast_next_sdf(
31     history_sdfs: List[np.ndarray],
32     history_times_h: np.ndarray,
33     mean_field: np.ndarray,
34     Vt: np.ndarray,
35     z_mean: np.ndarray,
36     z_std: np.ndarray,
37     k_llm: int,
38     meta_df=None,
39     horizon: int = 1,
40     rollout_nonce: Optional[float] = None,
41     caps: Optional[Dict] = None,
42     knn_lib: Optional[Dict] = None,
43 ) -> Tuple[np.ndarray, Dict]:
44     """Single-step SDF forecast. Returns (sdf_next, debug_dict)."""

```

```

45
46 def rollout_mc(
47     full_gt_imgs: List[np.ndarray],
48     times_h_all: np.ndarray,
49     start_t: int,
50     mean_field: np.ndarray,
51     Vt: np.ndarray,
52     z_mean: np.ndarray,
53     z_std: np.ndarray,
54     k_llm: int,
55     caps: Optional[Dict] = None,
56     meta_df=None,
57     knn_lib: Optional[Dict] = None,
58     n_rollouts: int = 8,
59     verbose: bool = True,
60 ) -> Tuple[List[Dict], List[np.ndarray]]:
61     """Run R autoregressive rollouts from start_t.
62     Returns (records, final_preds)."""

```

### 10.8.2 Single-Step Pipeline (forecast\_next\_sdf)

The function implements the complete single-step pipeline:

1. Project history SDFs into normalised PCA latent space.
2. Compute kNN prior delta  $\delta_{\text{prior}}$  and uncertainty  $\sigma$ .
3. Compute residual caps:  $\|\delta_{\text{LLM}}\|_2 \leq \text{RESIDUAL\_NORM\_FRAC} \cdot c_{\text{mag}}$  and per-component cap  $\max(0.25, 3 \cdot \text{median}(\sigma))$ .
4. Call LLM for residual correction.
5. Combine:  $\delta = \delta_{\text{prior}} + \alpha \cdot \delta_{\text{LLM}}$ .
6. Apply training caps (Section 5.1).
7. Apply velocity-relative cap (Equation 23).
8. Update latent:  $\tilde{\mathbf{z}}_{t+1} = \tilde{\mathbf{z}}_t + \delta$ .
9. Reconstruct SDF from updated physical-space latent.
10. Smooth SDF, enforce monotonic shrinkage (Section 5.4).

### 10.8.3 Dependencies

`numpy, data_loading, knn, llm_interface, metrics, pca, sdf_utils, config.`

## 10.9 autotune.py — PCA Auto-Tuning

**Purpose:** Performs a downsampled grid search over PCA hyperparameters ( $n_{\text{train}}, K$ ) to automatically select values that meet quality thresholds while ensuring subspace stability.

### 10.9.1 Key Functions

```

1 def preload_autotune_fields(
2     slice_ids: List[int],
3     timesteps: List[int],
4     downsample_factor: int = 2,
5 ) -> None:

```

```

6     """Eagerly load fields into cache for fast grid search."""
7
8 def fit_pca_on_indices(
9     indices: List[int],
10    k_fit_max: int,
11    timesteps: List[int],
12    downsample_factor: int = 2,
13    seed: int = 0,
14    use_cache: bool = True,
15) -> Tuple[np.ndarray, np.ndarray, Tuple[int, int]]:
16     """Fit PCA basis on downsampled fields."""
17
18 def oracle_score_curve(
19     val_indices: List[int],
20     mean_field: np.ndarray,
21     Vt: np.ndarray,
22     K_grid: List[int],
23     timesteps: List[int],
24     downsample_factor: int = 2,
25     use_cache: bool = True,
26) -> pd.DataFrame:
27     """Evaluate oracle reconstruction for each K.
28     Returns DataFrame with columns K, iou_mean, dice_mean."""
29
30 def subspace_distance_deg(
31     VtA: np.ndarray, VtB: np.ndarray, K: int
32) -> Tuple[float, float]:
33     """Mean and max principal angle (degrees) between
34     two K-dim subspaces."""
35
36 def auto_tune_pca(
37     train_max_index: int, seed: int = 0
38) -> Tuple[int, int, int]:
39     """Grid-search for (n_train, K_llm, K_PCS_MAX).
40     Returns (tuned_n_train, tuned_k_llm, tuned_k_recon)."""

```

### 10.9.2 Grid Configuration

The search grid is defined in `config.py`:

- `TUNE_TRAIN_GRID`: [400, 800] training slices.
- `TUNE_K_GRID`: [40, 60, 80, 120, 160] components.
- `TUNE_TIMESTEPS`: [0, 3, 6, 9] sampled time-steps.
- `TUNE_DOWNSAMPLE`: factor 2 spatial downsampling.
- `TUNE_VAL_SLICES`: 25 held-out validation slices.

### 10.9.3 Selection Logic

1. For each  $n_{\text{train}}$ , fit PCA on downsampled training fields.
2. Evaluate oracle reconstruction on validation set for each  $K$ .
3. Select smallest  $K_{\text{LLM}}$  meeting IoU/Dice/BF1 thresholds.
4. Select  $K_{\text{recon}}$  at diminishing-returns elbow.
5. Verify subspace stability via principal angles (Section 8.4).
6. Among all stable candidates, select the one with the smallest  $n_{\text{train}}$  and then the smallest  $K_{\text{LLM}}$ .

#### 10.9.4 Dependencies

`numpy, pandas, sklearn (optional), data_loading, metrics, pca, sdf_utils, config.`

### 10.10 ablation.py — Ablation Study Runner

**Purpose:** Compares pipeline variants by temporarily overriding configuration parameters and running the evaluation loop. Supports checkpoint/resume via `pickle` serialisation.

#### 10.10.1 Key Functions

```

1 def run_ablation_variant(
2     name: str,
3     valid_slices: List[int],
4     times_h_all: np.ndarray,
5     mean_field: np.ndarray,
6     Vt: np.ndarray,
7     z_mean: np.ndarray,
8     z_std: np.ndarray,
9     k_llm: int,
10    caps: Optional[Dict],
11    meta_df,
12    knn_lib: Optional[Dict],
13    n_rollouts: int,
14    llm_temperature: Optional[float] = None,
15    residual_scale: Optional[float] = None,
16    use_knn_guide: Optional[bool] = None,
17 ) -> Tuple[List[Dict], float]:
18     """Run one ablation variant. Returns (rows, runtime_s)."""
19
20 def run_all_ablations(
21     valid_slices: List[int],
22     times_h_all: np.ndarray,
23     mean_field: np.ndarray,
24     Vt: np.ndarray,
25     z_mean: np.ndarray,
26     z_std: np.ndarray,
27     k_llm: int,
28     caps: Optional[Dict],
29     meta_df,
30     knn_lib: Optional[Dict],
31 ) -> pd.DataFrame:
32     """Run all four ablation variants with checkpoint/resume."""

```

#### 10.10.2 Ablation Variants

Variant	Rollouts	Temperature	Residual scale	kNN
MC ensemble (baseline)	$R = 8$	0.60	0.70	yes
Single-shot	1	0.60	0.70	yes
Deterministic	1	0.00	0.70	yes
kNN-only (no LLM)	1	0.60	0.00	yes

Table 2: Ablation study variants.

#### 10.10.3 Checkpoint/Resume

After each variant completes, the accumulated rows are serialised to `ablation_partial.pkl`. On restart, completed variants (identified by name) are skipped. This ensures that long-running

ablation studies survive interruptions.

#### 10.10.4 Dependencies

`pandas, numpy, requests, data_loading, forecasting, config.`

### 10.11 plotting.py — Visualisation Routines

**Purpose:** Provides 11 publication-ready visualisation functions using only `matplotlib` (no `seaborn` dependency). All functions apply a consistent style via `rcParams`.

#### 10.11.1 Available Plots

1. `plot_mask_timeseries` — Raw mask sequence across all time-steps for a single slice.
2. `plot_area_vs_time` — Material pixel count versus degradation time.
3. `plot_explained_variance` — Cumulative PCA explained variance versus  $K$ .
4. `plot_pca_reconstructions` — Side-by-side ground truth versus PCA reconstructions at various  $K$  values.
5. `plot_eigenimages` — First  $n$  principal component images (“corrosion modes”).
6. `plot_metric_vs_horizon` — Mean  $\pm$  std of a metric (IoU, Dice, area error) versus forecast horizon.
7. `plot_ablation_comparison` — Multi-variant comparison (mean  $\pm$  std) on the same axes.
8. `plot_delta_over_baseline` —  $\Delta$ IoU gain of each variant over the kNN-only baseline.
9. `plot_gallery` — Side-by-side GT, mean prediction, XOR difference, and  $P(\text{material})$  uncertainty map.
10. `plot_reliability_diagram` — Calibration curve with ECE annotation.
11. `plot_risk_coverage` — Risk (mean pixel error) versus coverage for uncertainty utility assessment.

#### 10.11.2 Dependencies

`matplotlib, numpy, pandas.`

### 10.12 main.py — Pipeline Orchestrator

**Purpose:** Serves as the single entry point that orchestrates the entire pipeline from data loading to final plots.

#### 10.12.1 Execution Sequence

1. **Dataset discovery:** Count total slices; compute 80/20 train/test split index.
2. **Auto-tune PCA** (Section 8): if enabled, run grid search and update `cfg.NUM_TRAIN_SLICES`, `cfg.K_LLM`, `cfg.K_PCS_MAX`.
3. **Fit global PCA:** call `fit_global_pca` with tuned parameters.
4. **Load metadata:** degradation times from CSV.

5. **Compute training caps:** P95/P99 statistics on training deltas.
6. **Build kNN library:** feature–delta pairs from training sequences.
7. **Select test slices:** randomly sample NUM\_TEST\_SLICES from the test partition with valid data.
8. **Run MC rollouts:** for each test slice and each start\_t, perform NUM\_ROLLOUTS autoregressive forecasts. Collect per-step metrics and final predictions.
9. **Generate plots:** metric-vs-horizon, gallery, ablation comparison.
10. **Run ablation study:** four variants with checkpoint/resume.
11. **Uncertainty diagnostics:** compute Brier, NLL, ECE; generate reliability diagram and risk–coverage curve.

### 10.12.2 Usage

```
export OPENROUTER_API_KEY="sk-or-v1-your-key-here"
python -m corrosion_forecast.main
```

### 10.12.3 Dependencies

All other modules in the package.

## 11 Data Flow Diagram

Figure 1 shows the end-to-end data flow through the pipeline. Each box represents a transformation stage, and arrows indicate data dependencies.

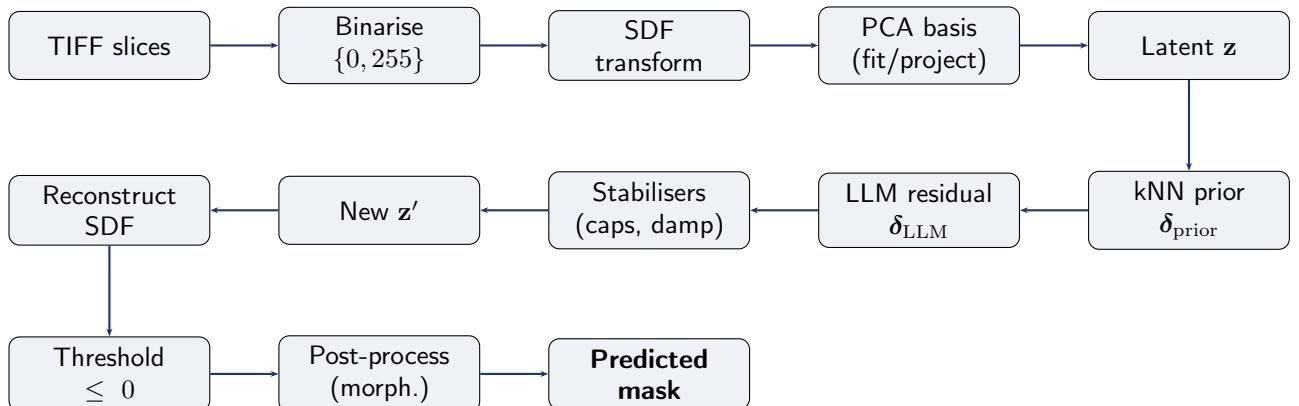


Figure 1: End-to-end data flow of the corrosion forecasting pipeline.

### Textual summary of the data flow:

TIFF slices → binarise → SDF → PCA basis → latent  $z$  → kNN prior → LLM residual → stabilise delta → new  $z'$  → reconstruct SDF → threshold → post-process → predicted mask.

Variable	Description
OPENROUTER_API_KEY	API key for the OpenRouter LLM endpoint. <b>Required</b> for any run that involves LLM calls.
CORROSION_BASE_DIR	Override the path to the <code>Corrosion_Masks</code> dataset directory. Defaults to <code>../Corrosion_Masks</code> relative to the package.

Table 3: Environment variables used by the pipeline.

## 12 Configuration and Reproducibility

### 12.1 Environment Variables

### 12.2 Modifying Hyperparameters

All hyperparameters are defined as module-level constants in `config.py` (Section 10.1). To change a parameter for an experiment:

1. Edit the constant directly in `config.py`.
2. Alternatively, override it programmatically before calling `main()`:

```

1 from corrosion_forecast import config as cfg
2 cfg.K_LLM = 60
3 cfg.NUM_ROLLOUTS = 16
4 cfg.LLM_TEMPERATURE = 0.8
5
6 from corrosion_forecast.main import main
7 main()

```

### 12.3 Random Seeds

The pipeline uses a fixed seed (`seed=0` by default) for all stochastic operations: NumPy random, Python `random` module, and the `random_state` parameter of randomised SVD. The kNN library construction and PCA training slice sampling are both seeded. The LLM itself is inherently stochastic ( $\text{temperature} > 0$ ), which is the *intended* source of variability across Monte Carlo rollouts.

### 12.4 Dataset Layout

The expected directory structure under `BASE_DIR` is:

```

Corrosion_Masks/
|--- metadata.csv
|--- Processed_0/
|   |--- 0.tif
|   |--- 1.tif
|   |--- ...
|--- Processed_1/
|   |--- 0.tif
|   |--- ...
|--- ...
`--- Processed_9/
    |--- 0.tif
    |--- ...

```

Each `Processed_t/` directory contains one TIFF file per slice index, and the integer  $t$  indexes the degradation time-step (0 through 9). The file `metadata.csv` contains a column `Degradation Time (h)` with the physical time in hours for each time-step.

## 12.5 Reproducibility Checklist

1. Set `OPENROUTER_API_KEY` and, if necessary, `CORROSION_BASE_DIR`.
2. Install dependencies: `pip install -r corrosion_forecast/requirements.txt`.
3. Run: `python -m corrosion_forecast.main`.
4. Results will be printed to stdout; plots will be displayed interactively or saved if `plt.savefig` calls are added.
5. The ablation checkpoint file (`ablation_partial.pkl`) enables resuming interrupted runs.

## References

- [1] N. Halko, P. G. Martinsson, and J. A. Tropp, “Finding structure with randomness: Probabilistic algorithms for constructing approximate matrix decompositions,” *SIAM Review*, vol. 53, no. 2, pp. 217–288, 2011.
- [2] F. Pedregosa *et al.*, “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [3] OpenRouter, “Unified API for LLMs,” <https://openrouter.ai/>, 2024.
- [4] L. R. Dice, “Measures of the amount of ecologic association between species,” *Ecology*, vol. 26, no. 3, pp. 297–302, 1945.
- [5] M. P. Naeini, G. Cooper, and M. Hauskrecht, “Obtaining well calibrated probabilities using Bayesian binning,” *Proc. AAAI Conference on Artificial Intelligence*, 2015.
- [6] G. Csurka, D. Larlus, and F. Perronnin, “What is a good evaluation measure for semantic segmentation?” *Proc. BMVC*, 2013.
- [7] J. A. Sethian, *Level Set Methods and Fast Marching Methods*, Cambridge University Press, 2nd edition, 1999.
- [8] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” *Proc. ICML*, 2017.